

The Importance of Monitoring

1. Overview

Monitoring is one of the **most critical pillars** of production systems.

A system without monitoring is like:

Flying an aircraft with no dashboard.

Even if architecture is excellent, without visibility you won't know:

- Failures
- Slowdowns
- Resource exhaustion
- Security anomalies
- Customer pain

2. Definition

Monitoring is the continuous observation of system health, performance, reliability, and user experience using metrics, logs, traces, and alerts.

3. Why Monitoring Matters

Monitoring helps detect:

Technical Problems

- Server crashes

- High CPU / memory
- Database slow queries
- Queue backlog
- Cache failures

Performance Problems

- Increased latency
- Timeout spikes
- Throughput drops

User Experience Problems

- Checkout failures
- Login issues
- Video buffering
- Push notifications delayed

Security Problems

- Suspicious traffic
- Brute-force login attempts
- Data access anomalies

4. Golden Rule

If you cannot measure it, you cannot improve it.

5. Customer Experience is the Real Metric

System may look healthy internally, but users may suffer.

Examples:

- APIs return 200 but wrong data
- UI loads slowly due to CDN issue
- Payment success internally but user sees failure screen

So monitoring should include:

External Signals

- Synthetic checks
- Real User Monitoring (RUM)
- Conversion funnels

6. Core Monitoring Pillars

A. Metrics

Numeric time-series data.

Examples:

- Requests/sec
- Error %
- CPU usage
- Memory usage
- DB connections
- Cache hit rate

B. Logs

Detailed event records.

Examples:

None

User 123 login failed

Order 456 payment timeout

C. Traces

Track request across microservices.

None

Client → API Gateway → Auth → Order → Payment

Useful for latency bottlenecks.

D. Alerts

Automatic notifications when thresholds breached.

Examples:

- Error rate > 5%
- P99 latency > 800ms
- Disk > 90%

7. Incident Priority Levels

You mentioned urgency - very important.

P0 (Critical)

- Entire service down
- Payment broken
- Massive outage

Immediate action.

P1 (High)

- Major degradation
- Region down
- Slow checkout

P2 (Medium)

- Non-critical feature broken

P3 (Low)

- Cosmetic issue
- Minor metric anomaly

8. Monitoring Dependencies

If your service is used by others:

None

Payments Service ← Order Service ← Checkout

When checkout breaks, they may blame payments.

You need monitoring to quickly prove:

- Was your service healthy?
- Was latency elevated?
- Any errors?

9. What to Monitor

Application Layer

- QPS
- Success/error rate
- Latency P50/P95/P99

Infrastructure Layer

- CPU
- Memory
- Disk
- Network

Database Layer

- Query latency
- Locks
- Replication lag
- Connections

Cache Layer

- Hit ratio
- Evictions
- Memory usage

Queue Layer

- Consumer lag
- Retry count
- DLQ growth

Business Layer

- Orders/min
- Revenue/hour
- Signup success %
- Churn rate

10. SLI / SLA / SLO

SLI

Measured metric.

Example:

- 99.9% successful requests

SLO

Target objective.

Example:

- P99 latency < 300ms

SLA

Contract with customer.

Example:

- 99.95% uptime

11. Example Dashboard

None

Traffic: 15k RPS

Error Rate: 0.3%

P99 Latency: 220ms

CPU: 62%

DB Lag: 0 sec

12. Good Monitoring Stack Examples

- Prometheus + Grafana
- Datadog
- New Relic
- ELK Stack
- OpenTelemetry

13. Interview Script

Say:

“I’d add observability from day one: metrics, logs, traces, dashboards, and alerting. I’d monitor not only infrastructure but also user-facing KPIs like checkout success and P99 latency.”

Strong answer.

14. Final Summary

- Monitoring provides production visibility
- Detects failures before users complain
- Includes metrics, logs, traces, alerts
- Monitor technical + business KPIs
- Critical for debugging dependencies

15. Memorable Line

Hope is not a monitoring strategy.